

Chapter 4

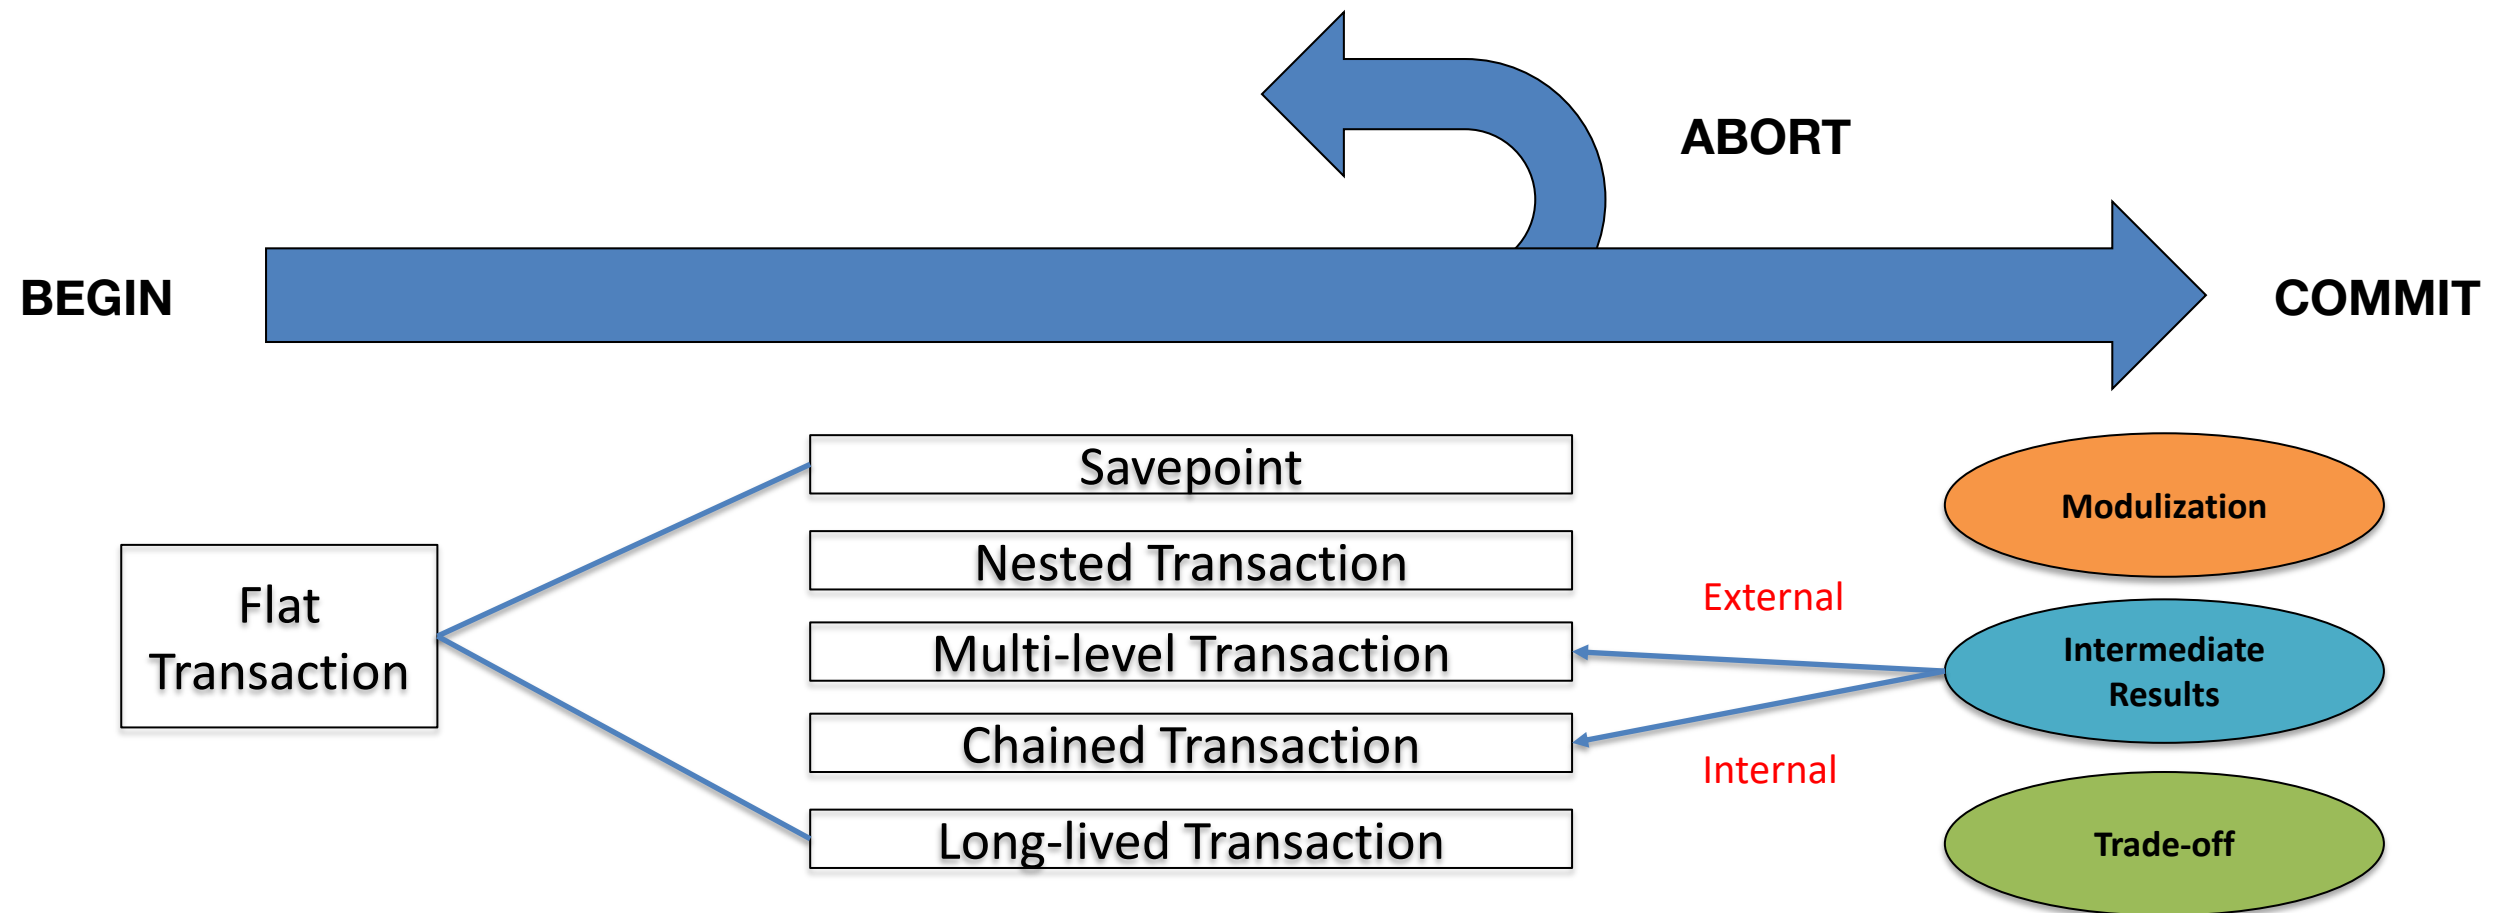
Transaction Models

Jong-Hyeok Park
akindo19@gmail.com



Introduction

- **Atomicity** is defined only from the perspective of the **caller** of the operation
- Trade-off : **Simplicity** of system design vs. **Price** for guaranteeing ACID



Atomic Actions

Disk Write as Atomic Action

1. Single disk write

- (1) Controller error (2) Power loss (3) Media fail

2. Read-after-write

- First issue write then re-read block from disk and compare the results with original block
- (1) Can't return to old version (2) Exceed write threshold (3) Undefined state @ crash

3. Duplexed write

- Write two places (like DWB strategy)
- Maintain version number and pick higher version at read

4. Logged write

- First write old block at different location (like RBJ SQLite – *before-image* logging)

Atomic Actions

Action Types

1. Unprotected
 - Lack all A, I, D properties
 - Must be controlled by application or embedded high-level protected action
2. Protected
 - Don't externalize result before they end thus, **ACID**
3. Real
 - Hard or impossible to reverse
 - Executed only if enclosing protected actions have reached a state will **not to be rollback**
 - **Testable** real actions are easy to recovery

Flat Transaction

- Simplest type of transaction
- Only **one layer** of control by the application *BEGIN ... COMMIT*
- **Atomicity**
 - All or Nothing
- **Consistency**
 - A transaction produce consistent results only
 - Commit itself is taken as the guarantee that the result is consistent.
- **Isolation**
 - Behave exactly as it would in single-user-mode
- **Durability**
 - Once transaction commit it must be reestablish its results after any type of subsequent failures

Flat Transaction

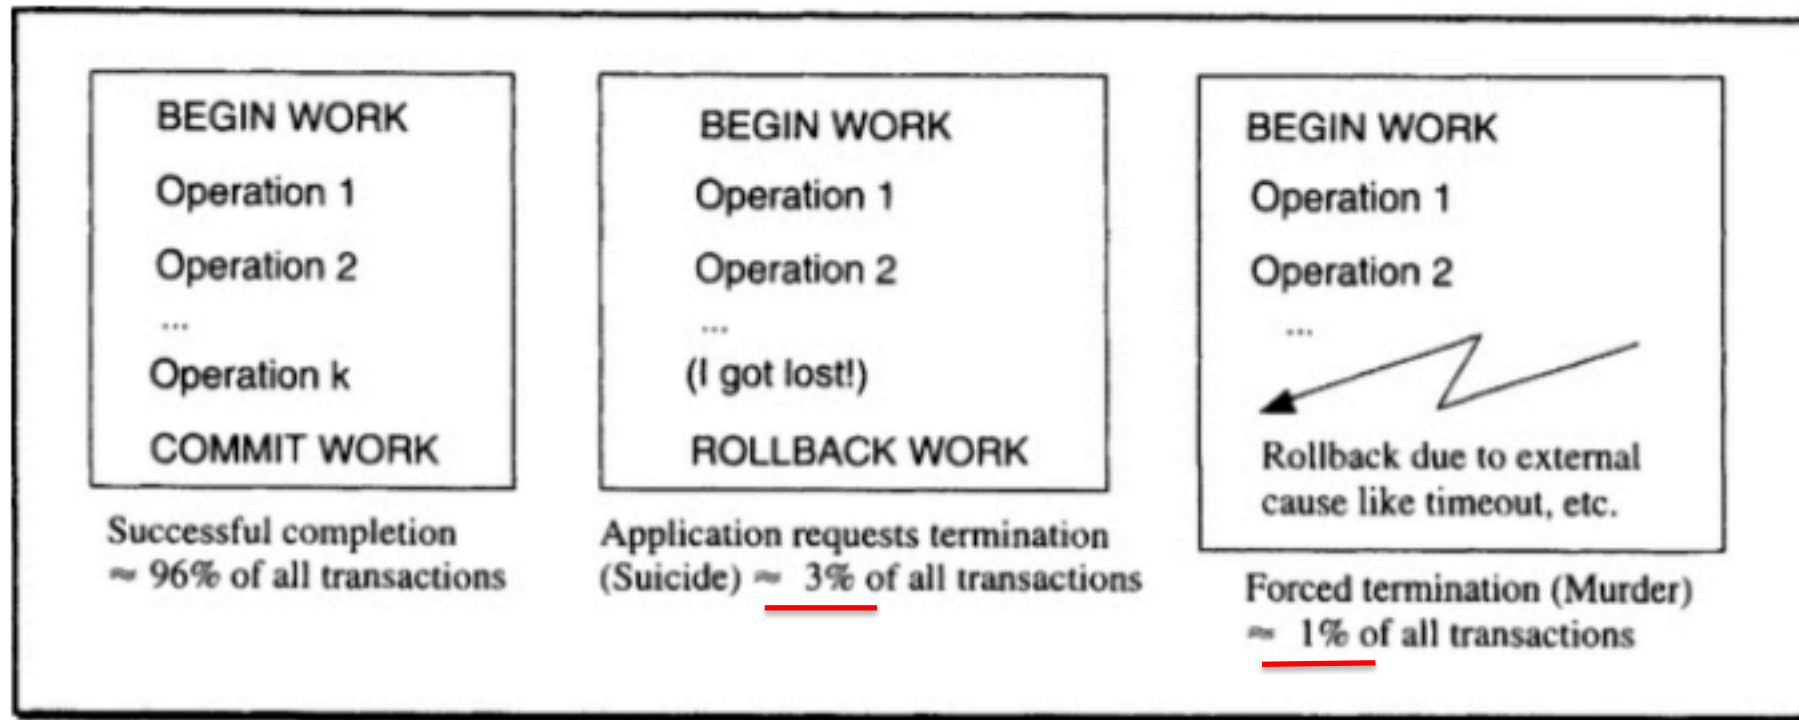


Figure 4.2

Flat Transaction

Limitations

- Selective rollback
- Bulk updates

BEGIN WORK

S1: book flight from San Francisco to Frankfurt
S2: book flight from Frankfurt to Milan, same day
S3: book flight from Milan to Pisa, same day

Problem: There is no way to get to Ripa from Pisa the same day, except in a rental car. But you do not want to drive at night.

Logging is dirty solution



Exceeding the capability of Flat transaction

```
ComputeInterest()
{
    real    interest_rate;    /* monthly interest rate */
    receive (interest_rate);  /* receive request to compute the accumulated */
                              /* interest and to modify all accounts */
                              /* accordingly */
    exec sql BEGIN WORK; /* start the transaction */
    exec sql UPDATE checking_accounts /*
        set account_balance = /*
            account_balance * (1+interest_rate);/* modify all accounts */
    send ("done"); /*
    exec sql COMMIT WORK; /*
    return; /*
};
```

Spheres of Control (SoC)

- **Process control**
 - Ensure that atomic process is **not modified by others** and **constrains dependencies**
- **Process atomicity**
 - Amount of processing one wishes to consider as having identity
- **Process commitment**
 - Commitment of the effects of process, even beyond the end of the process

Spheres of Control (SoC)

Case1. Constrain undesired dependency **before** they occur (House-buying example)

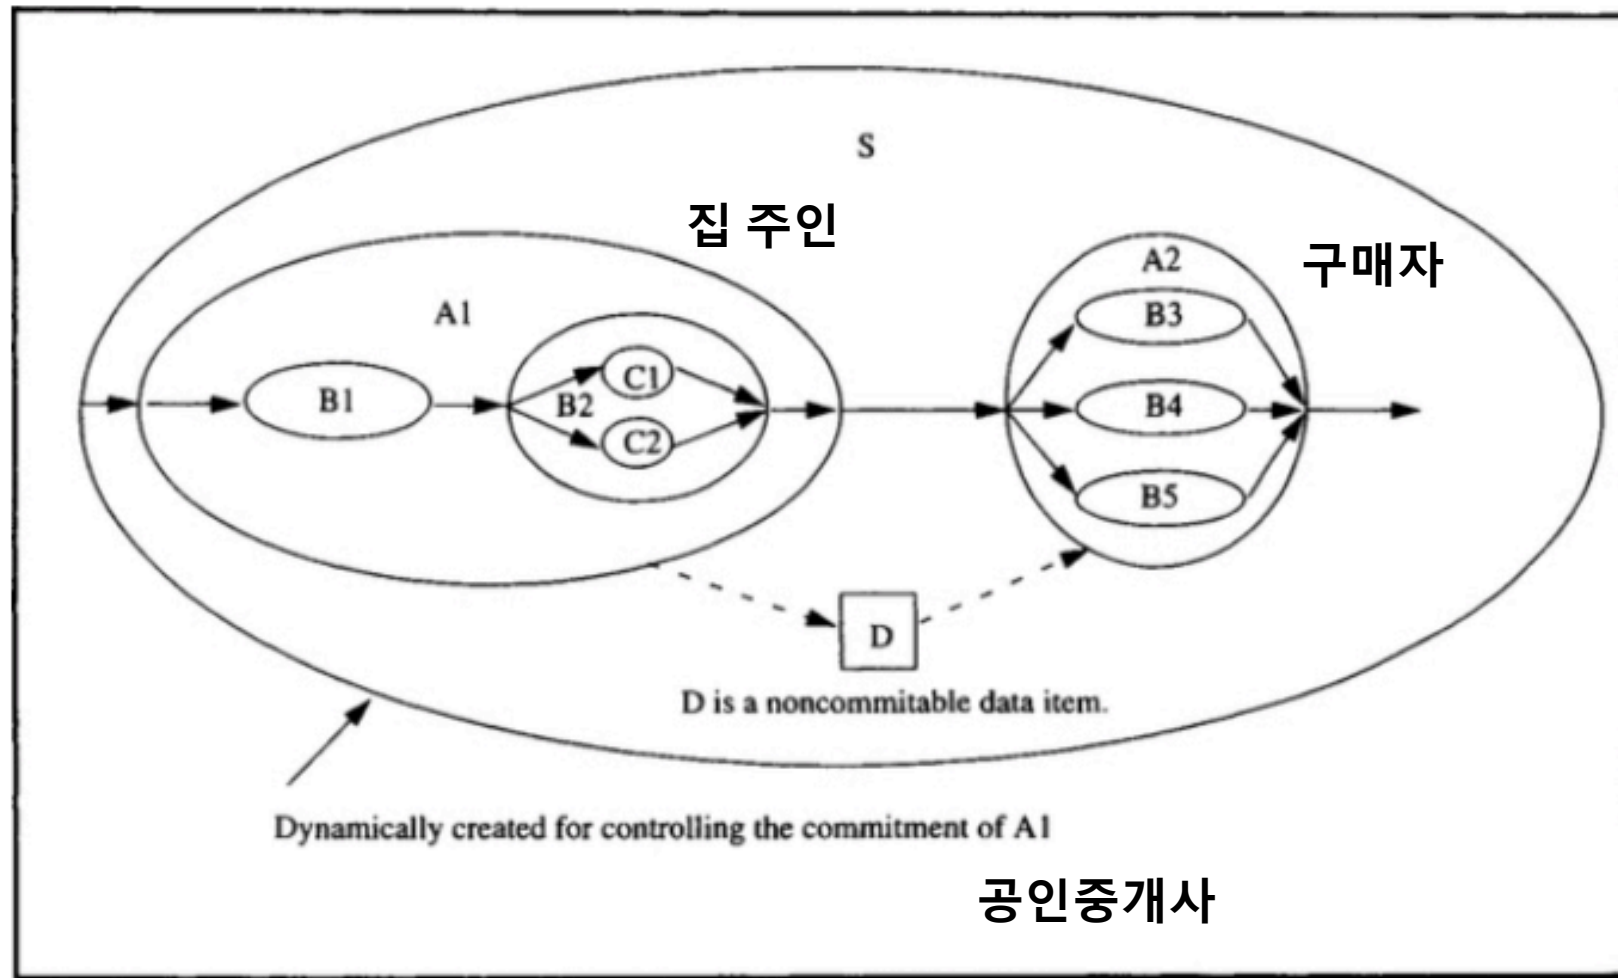
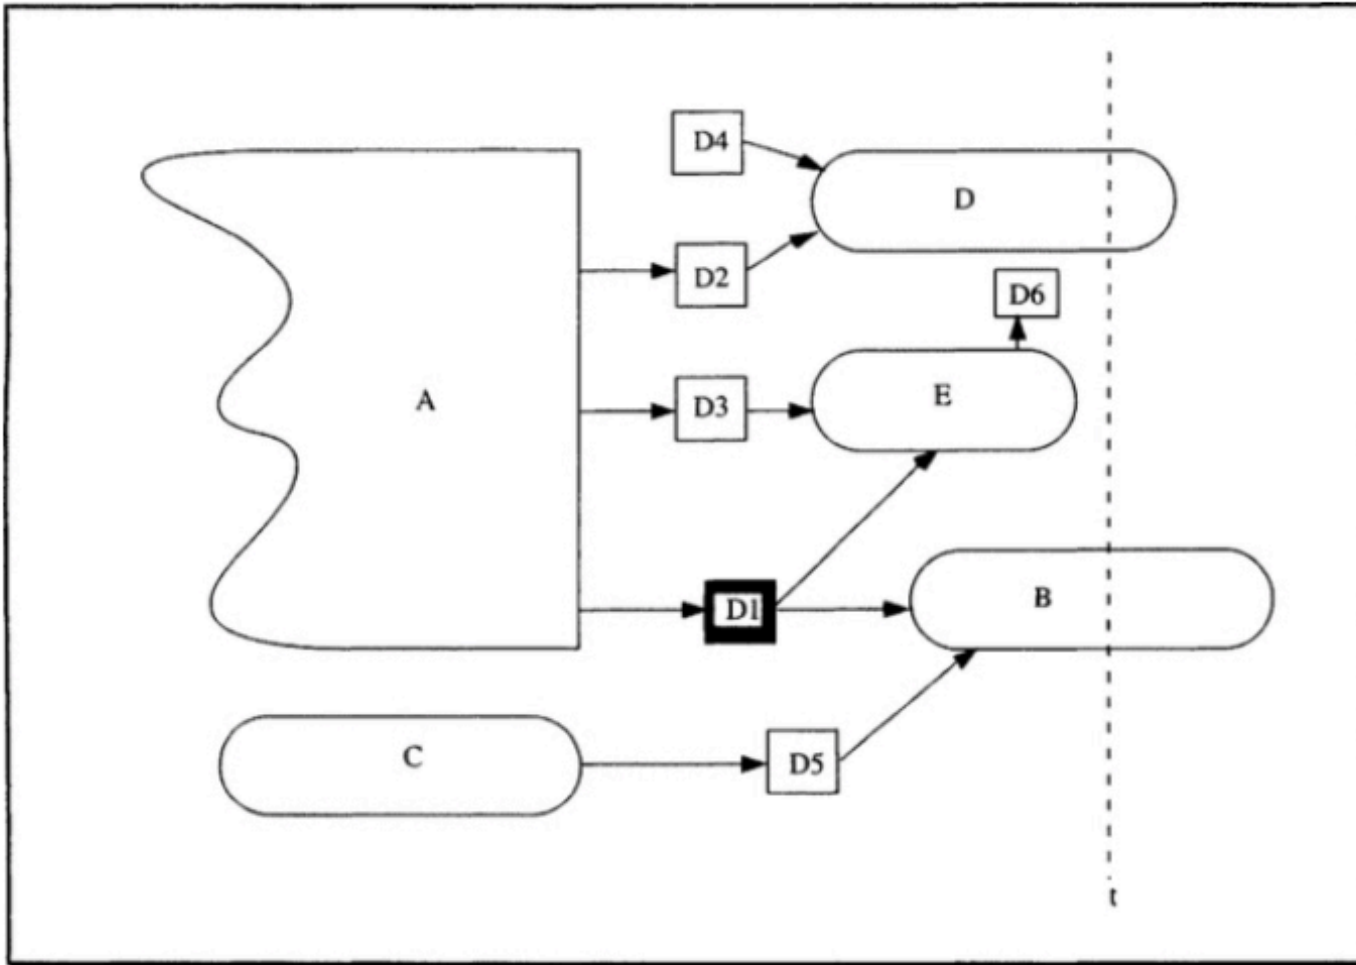


Figure 4.3

Spheres of Control (SoC)


Case2. Constrain undesired dependency **after** they occur



1. Trace back and find out culprit

Spheres of Control (SoC)

Case2. Constrain undesired dependency **after** they occur




1. Trace back and find out culprit

Extend backward SoC that invalid data item D1

Spheres of Control (SoC)

Case2. Constrain undesired dependency **after** they occur



1. Trace back and find out culprit

Extend backward SoC that invalid data item D1

2. Recovery SoC

- Create Dynamic dependencies to encompass all process affected by correction of an error
- All recovery step is application dependent

Notations

- Describe transaction model as *Finite State Machine* **NOT Appropriate !!!**

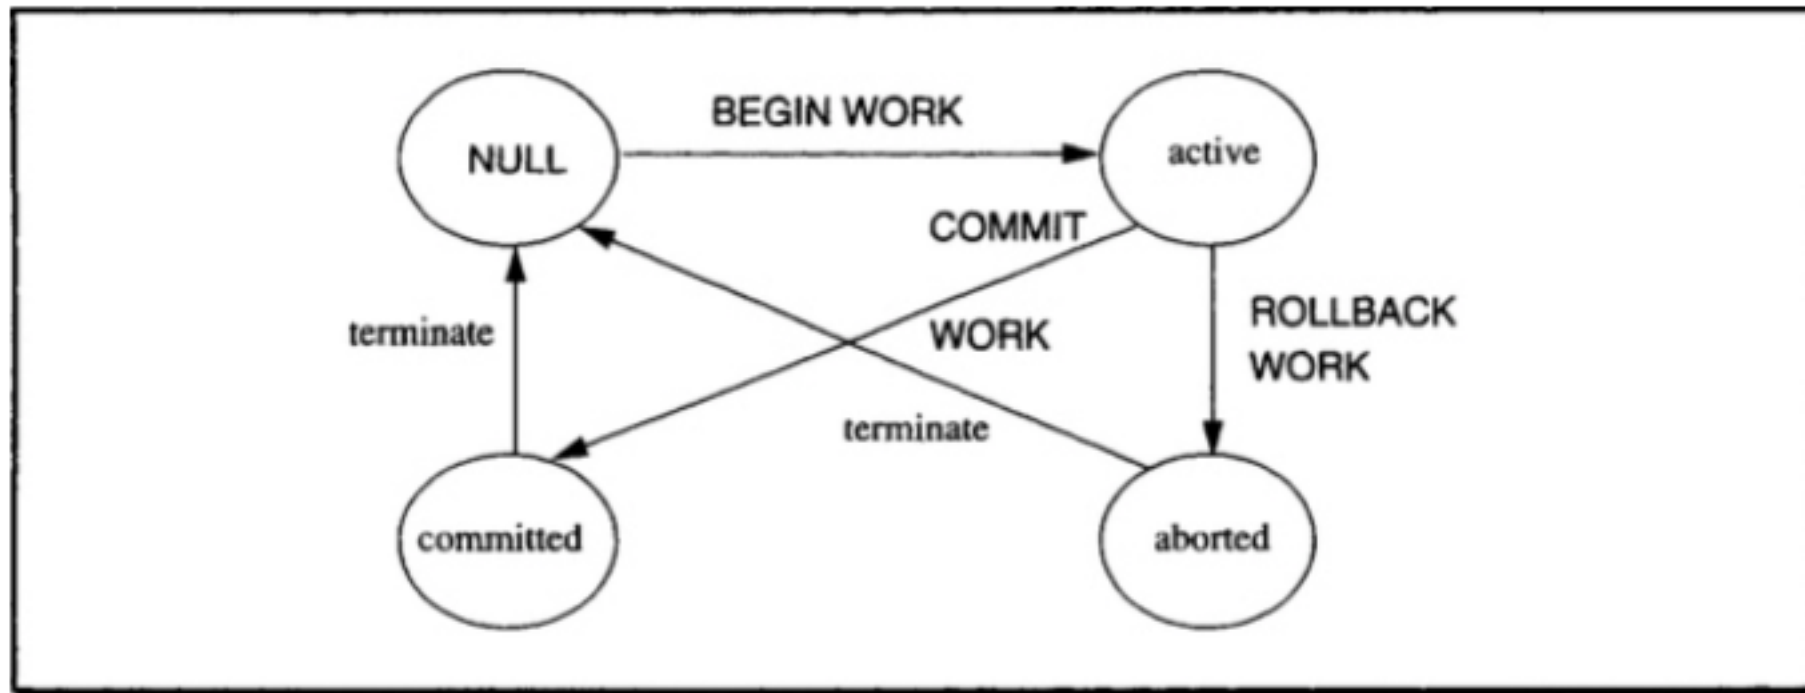
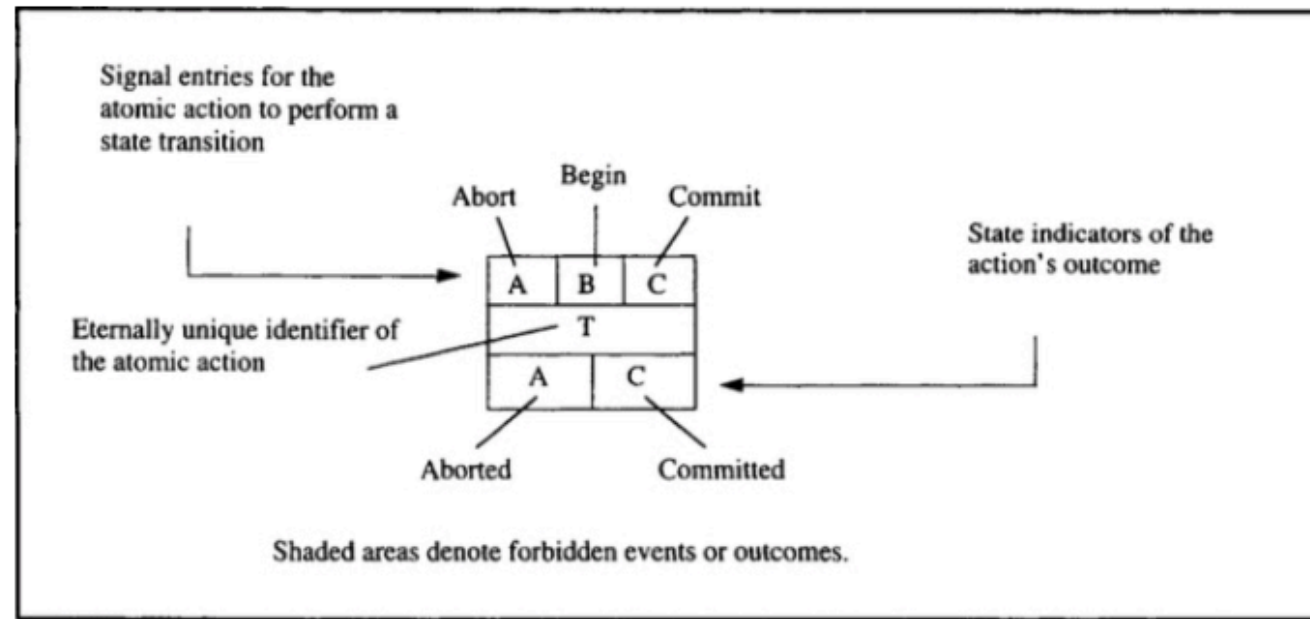


Figure 4.5

Notations

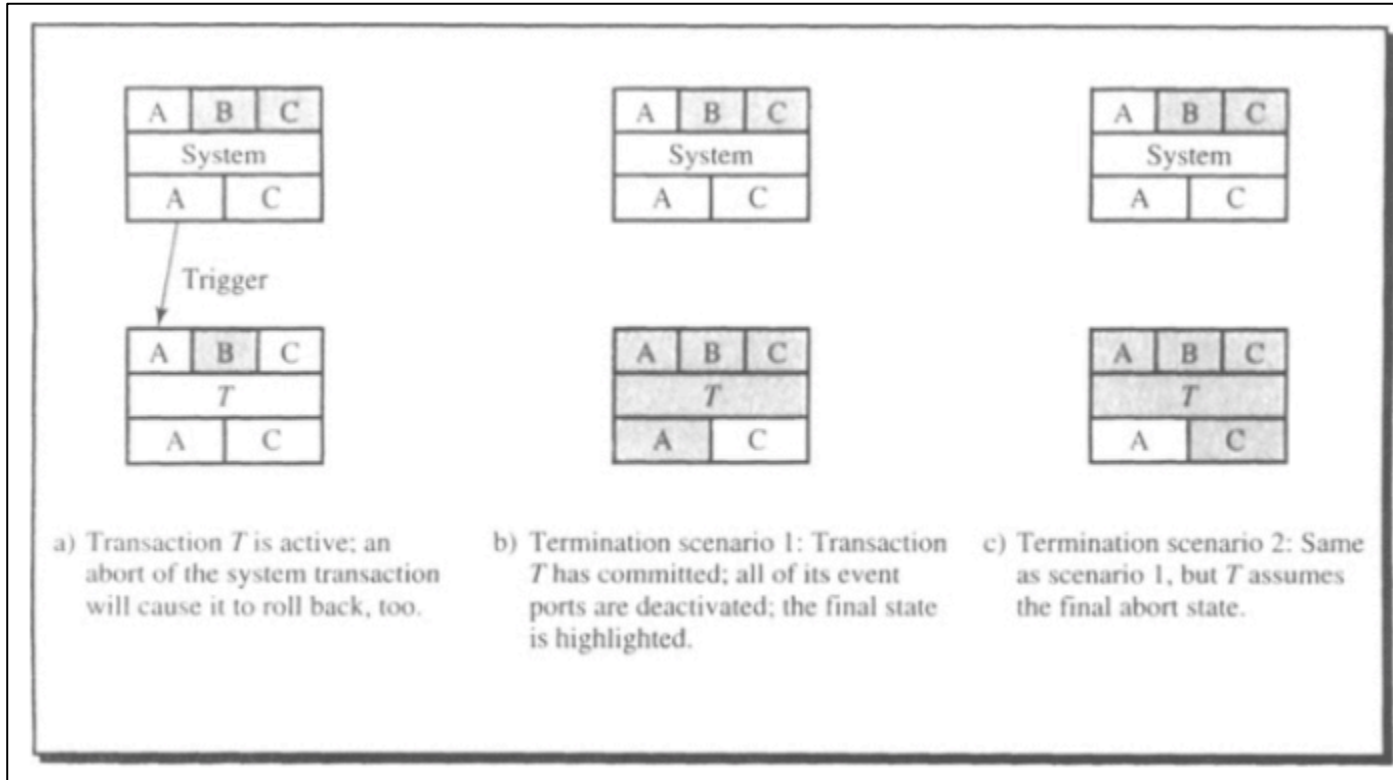
- Graphical Notations



- Rules

$\langle \text{rule identifier} \rangle : \langle \text{preconditions} \rangle \rightarrow \langle \text{rule modifier list} \rangle, \langle \text{signal list} \rangle, \langle \text{state transition} \rangle$

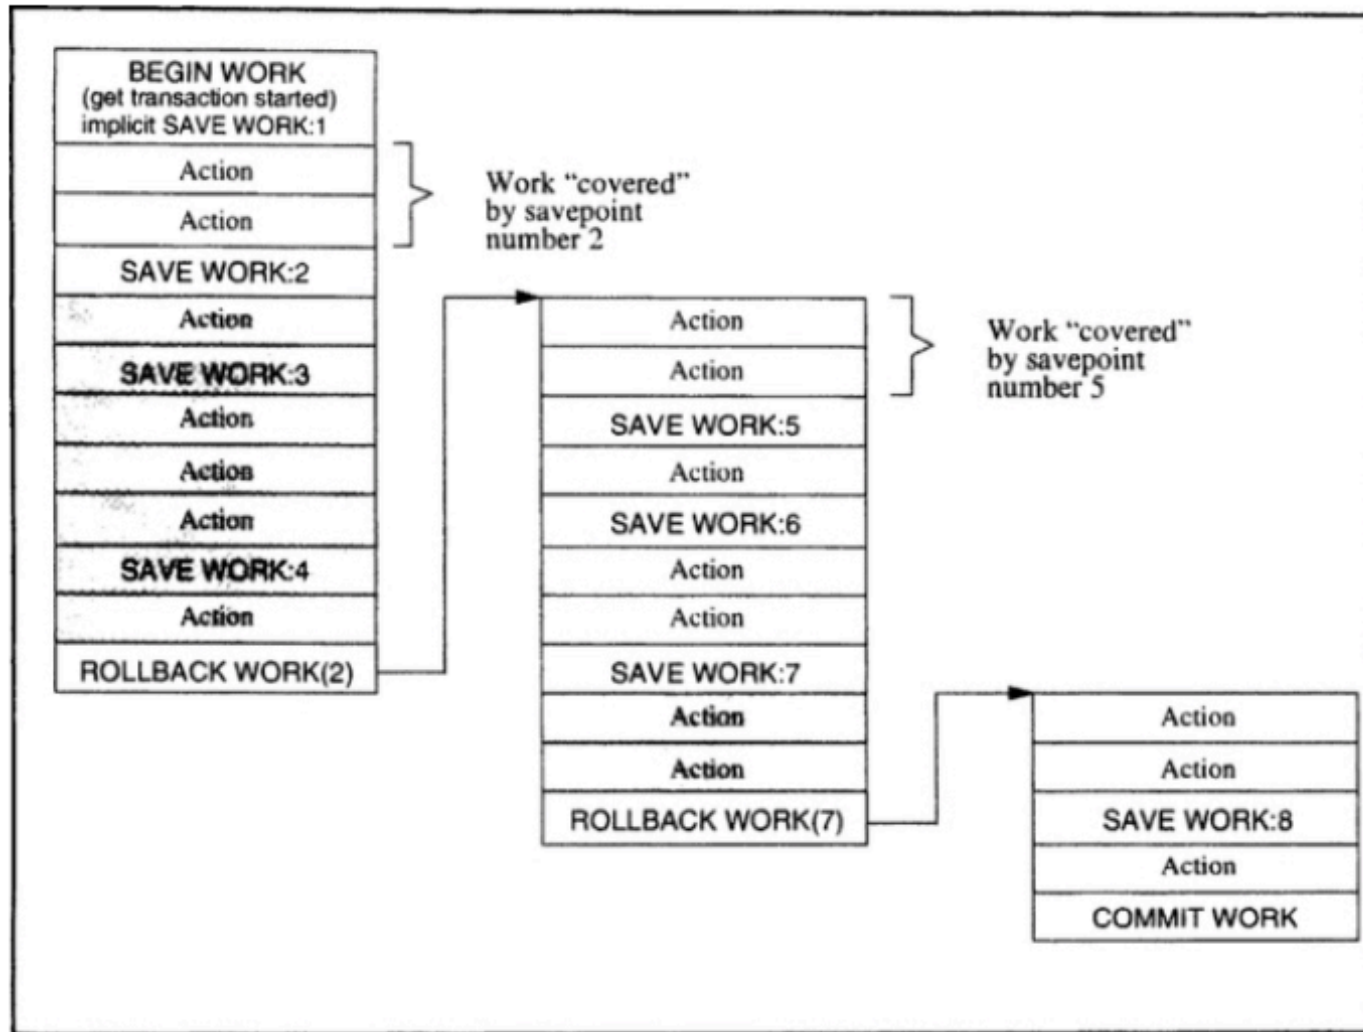
Notations



- **delete(X)** : All rules pertaining to X and all references to it are to be deleted.

$S_B(T) :$ $\rightarrow +(S_A(\text{system})|S_A(T)), \text{BEGIN WORK}$
 $S_A(T) :$ $\rightarrow (\text{delete}(S_B(T)), \text{delete}(S_C(T))), \text{ROLLBACK WORK}$
 $S_C(T) :$ $\rightarrow (\text{delete}(S_B(T)), \text{delete}(S_A(T))), \text{COMMIT WORK}$

Savepoint



1. Savepoint counter – monotonically increase
2. ROLLBACK does **NOT** affect Savepoint counter
3. When rolling back to Savepoint – TX **stay alive**
 - Keep its resources
 - Return initial state
4. When system is crashed, TX rolls back oldest atomic action


Savepoint

$S_B(S_n) :$	\rightarrow	„BEGIN WORK
$S_A(R) : (R < S_n)$	\rightarrow	, $S_A(S_n - 1)$, ROLLBACK WORK
$S_C(S_n) :$	\rightarrow	, $S_C(S_n - 1)$, COMMIT WORK
$S_S(S_n) :$	\rightarrow	+ ($S_A(S_n) S_A(S_n + 1)$), $S_B(S_n + 1)$,

- If the transaction rolls back from Savepoint-(N) to Savepoint-(N-k) then all atomic actions on way back are **aborted**
- **Persistent Savepoint**
 - Great reduction of work lost - Phoenix transaction
 - Not simple

Chained Transaction

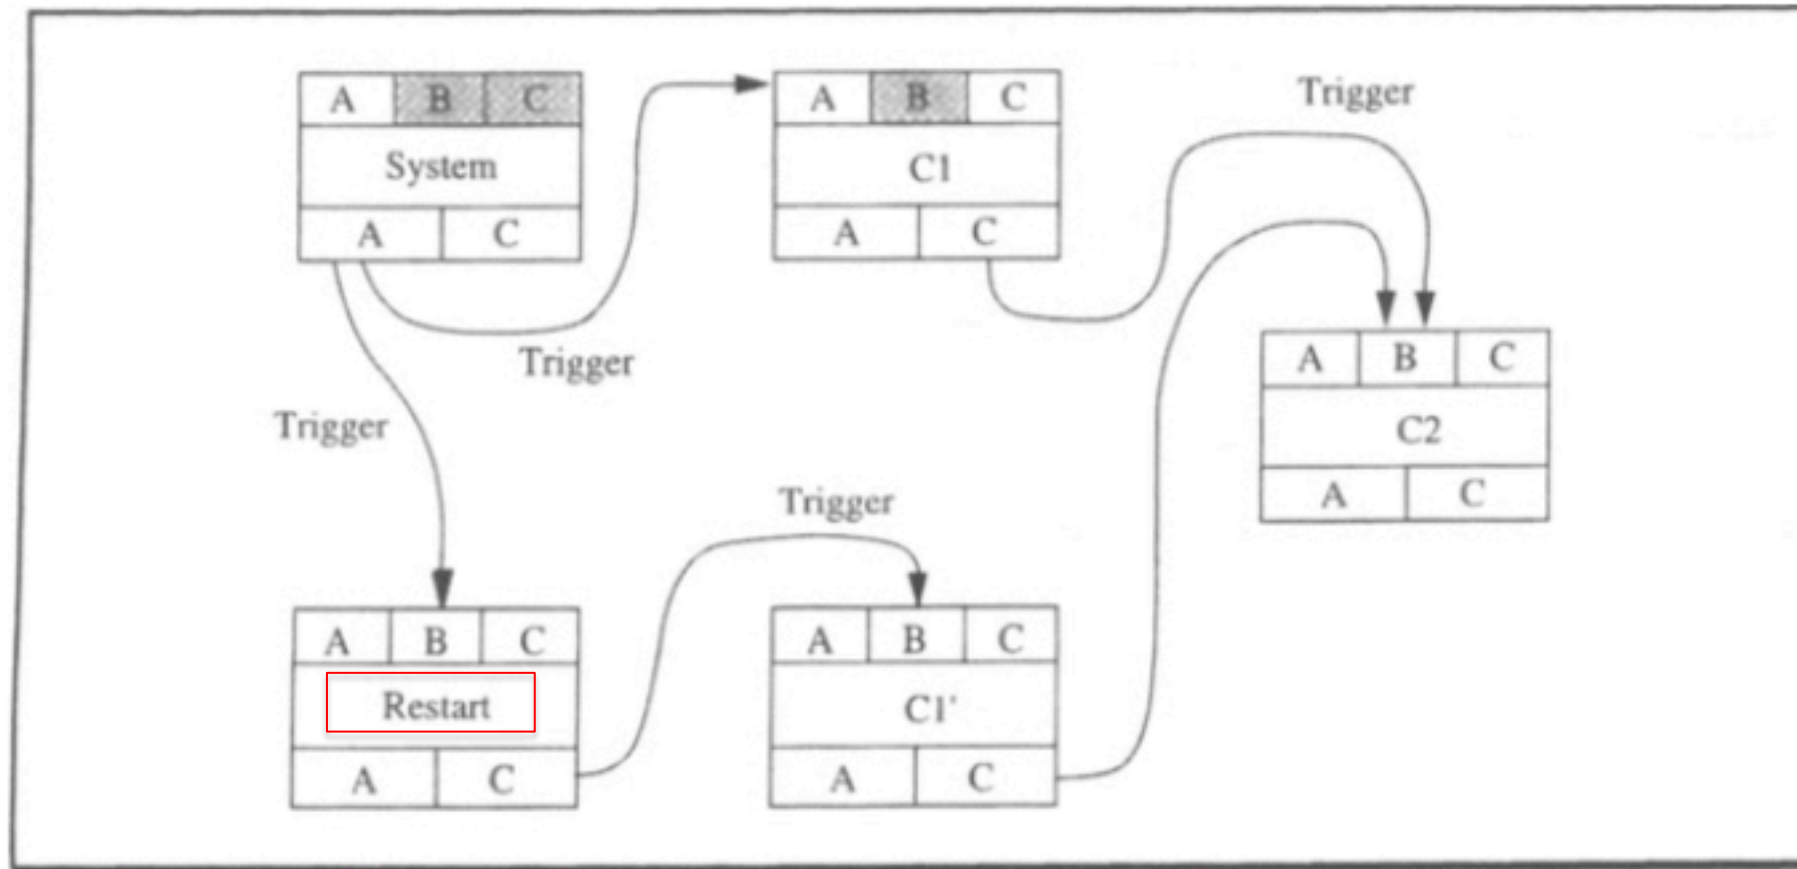
- CHAIN WORK = COMMIT + BEGIN



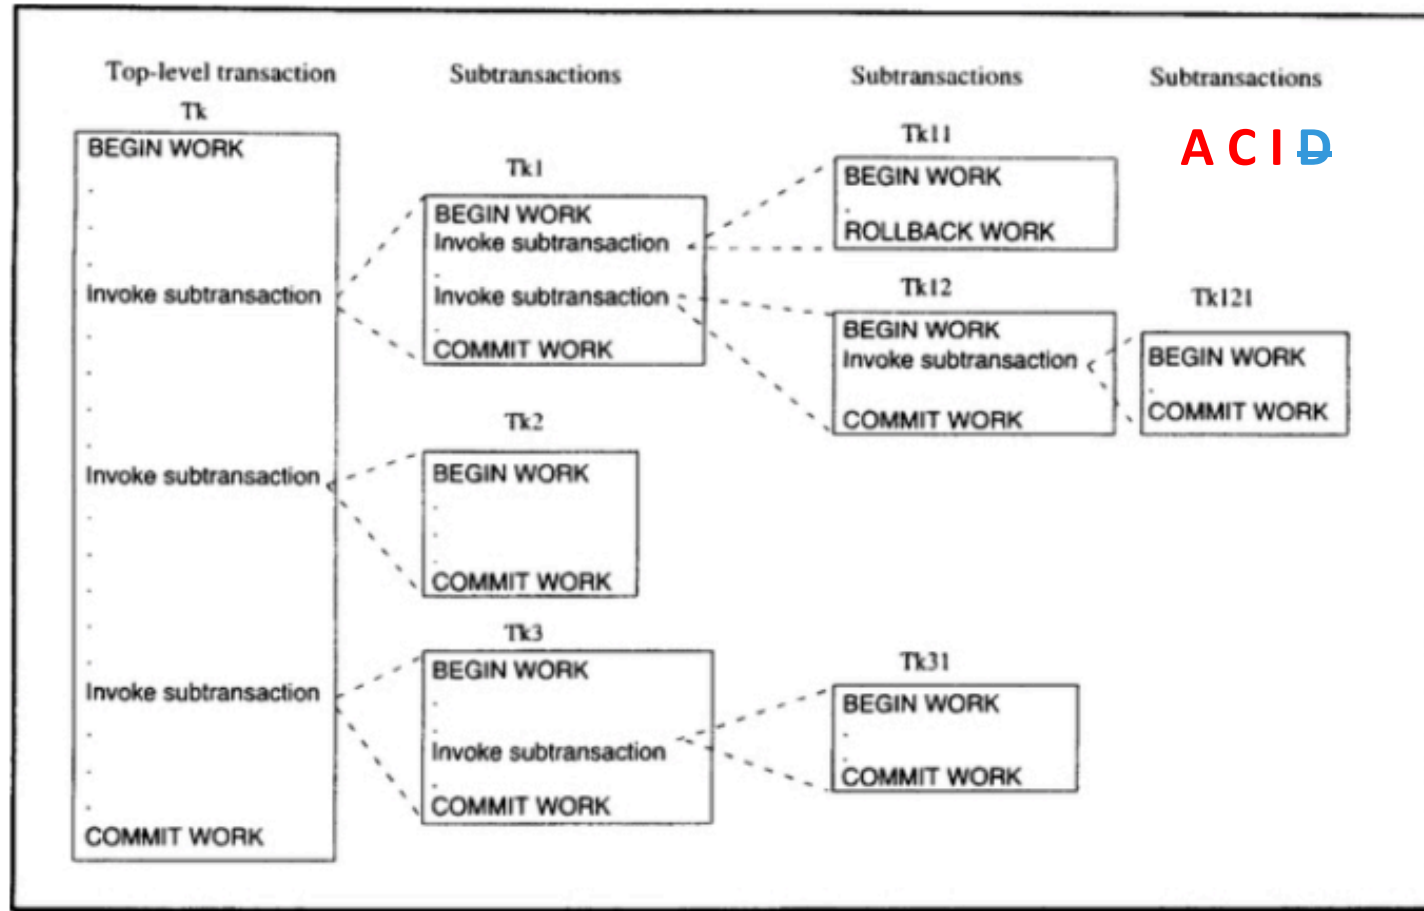
Vs. Savepoint

- Workflow structure**
 - Allow substructure (e.g. cursor)
- Commit versus Savepoint**
 - Restoring **previous savepoint only**
- Lock handling**
 - COMMIT frees all lock
- Work lost**
 - After crash, entire TXs are rolled back
- Restart handling**
 - Reestablish most recent commit

Chained Transaction



Nested Transaction

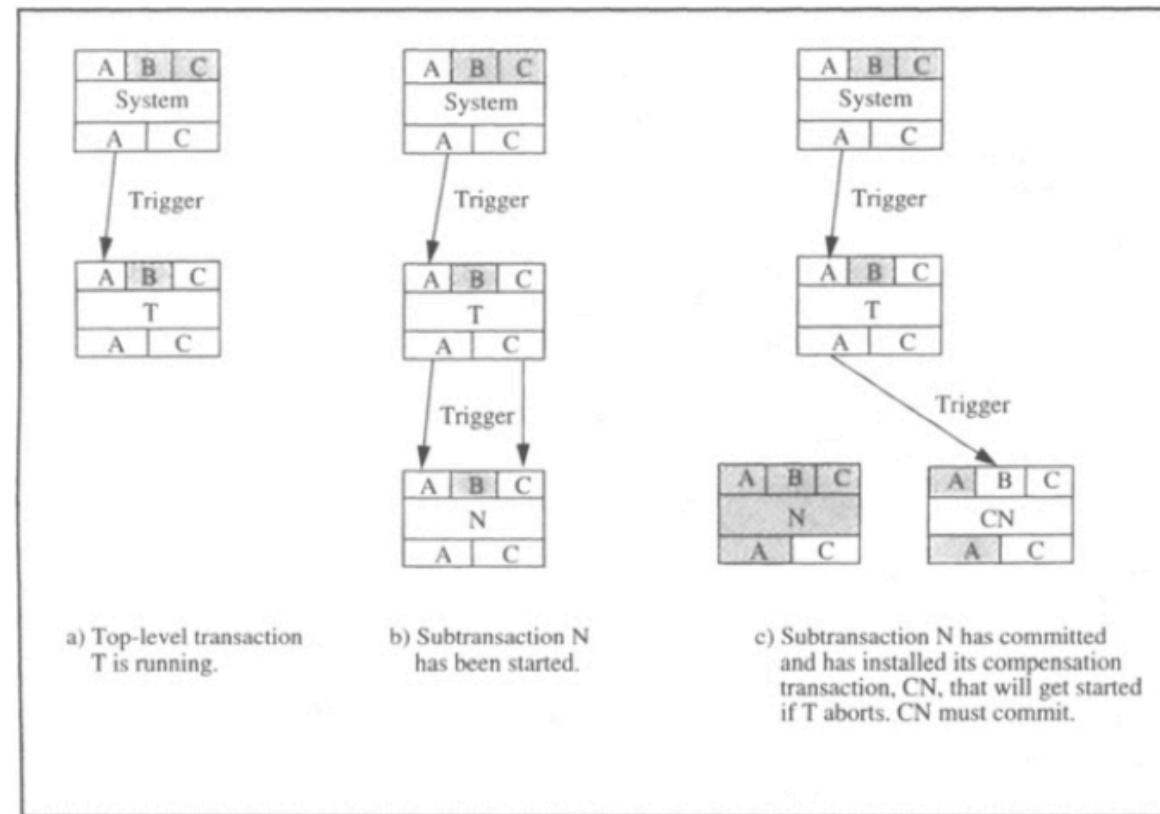


ACID

- Tree of transaction
- Leaf level transaction is *Flat transaction*
- **Commit Rule :**
Any subtransaction can finally commit **only if** the root transaction commits
- **Rollback Rule :**
Parent tx rolls back then all child txs must be rolled back.
- **Visibility Rule :**
Parent can see child upon it commits
Child can access Parent

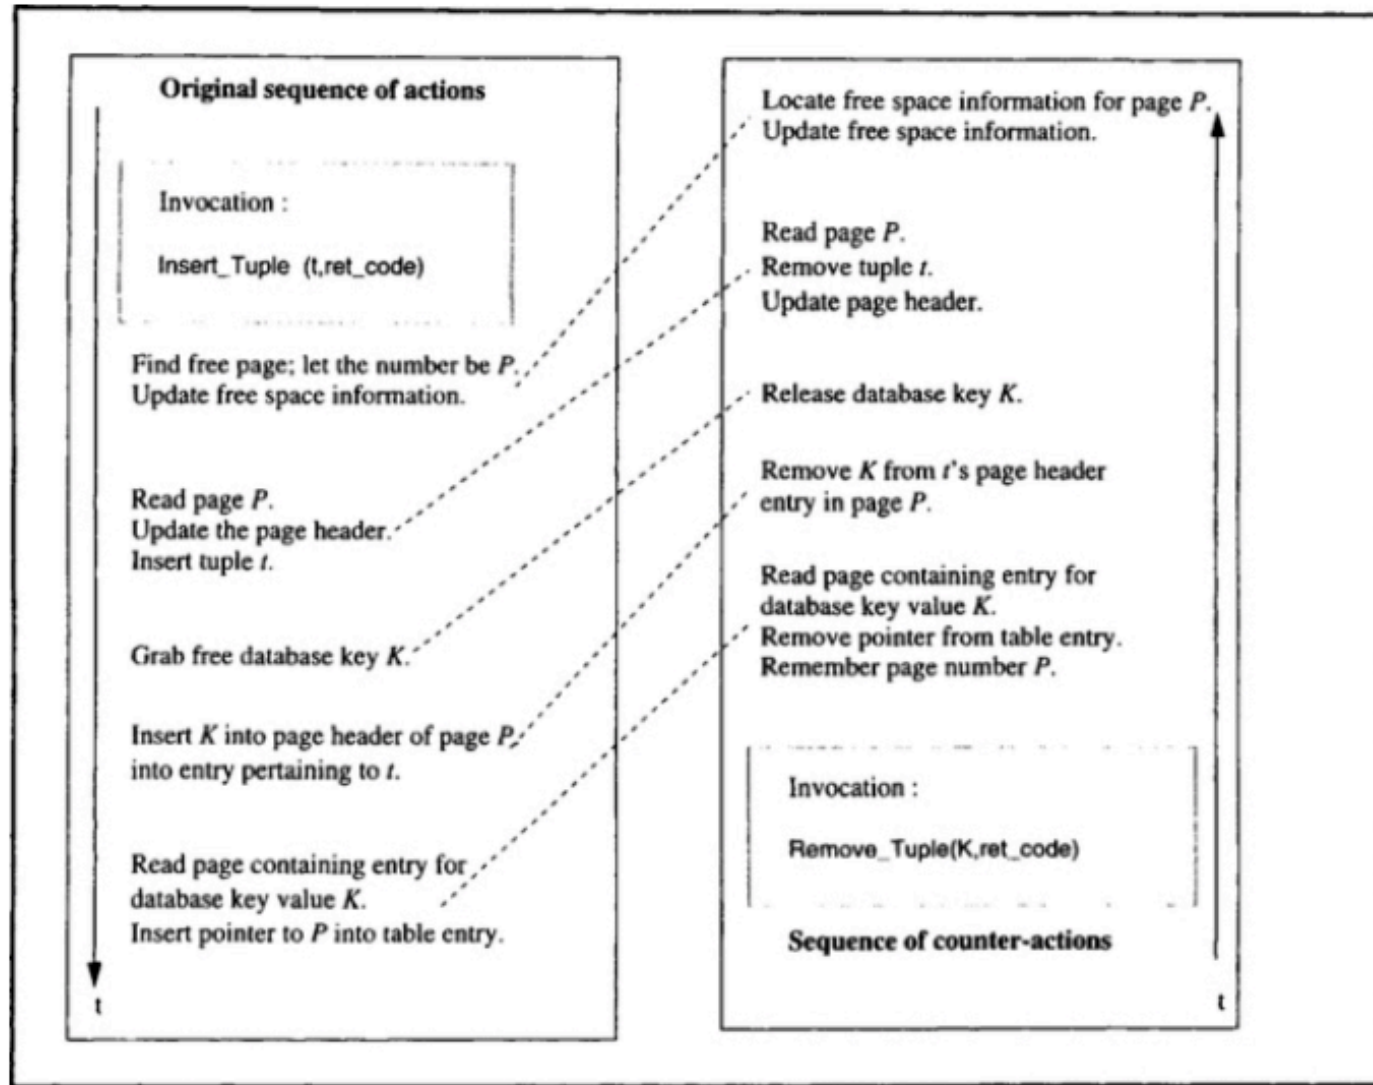
Multi-level Transaction

- General version of Nested Transaction
- Allow earlier commit (pre-commit) of sub-transactions
- Backout pre-committed result? - **Compensating Transaction**



Multi-level Transaction

- Compensating Transaction could be nested subtransactions




- Abstraction Hierarchy**
 - Entire system consists strict hierarchy of objects
- Layered abstraction**
 - The object layered n are completely implemented by using operation of layer $n-1$
- Discipline**
 - There are no shortcuts that allow layer n to access objects on a layer other than $n-1$.

Long-Lived Transaction

1. Minimize lost work – How to split up bulk transactions?
2. Recoverable computation – Temporarily stop computation without *commit* or *rollback*
3. Explicit control flow – System can control trx belonging to LLTs

SAGAs

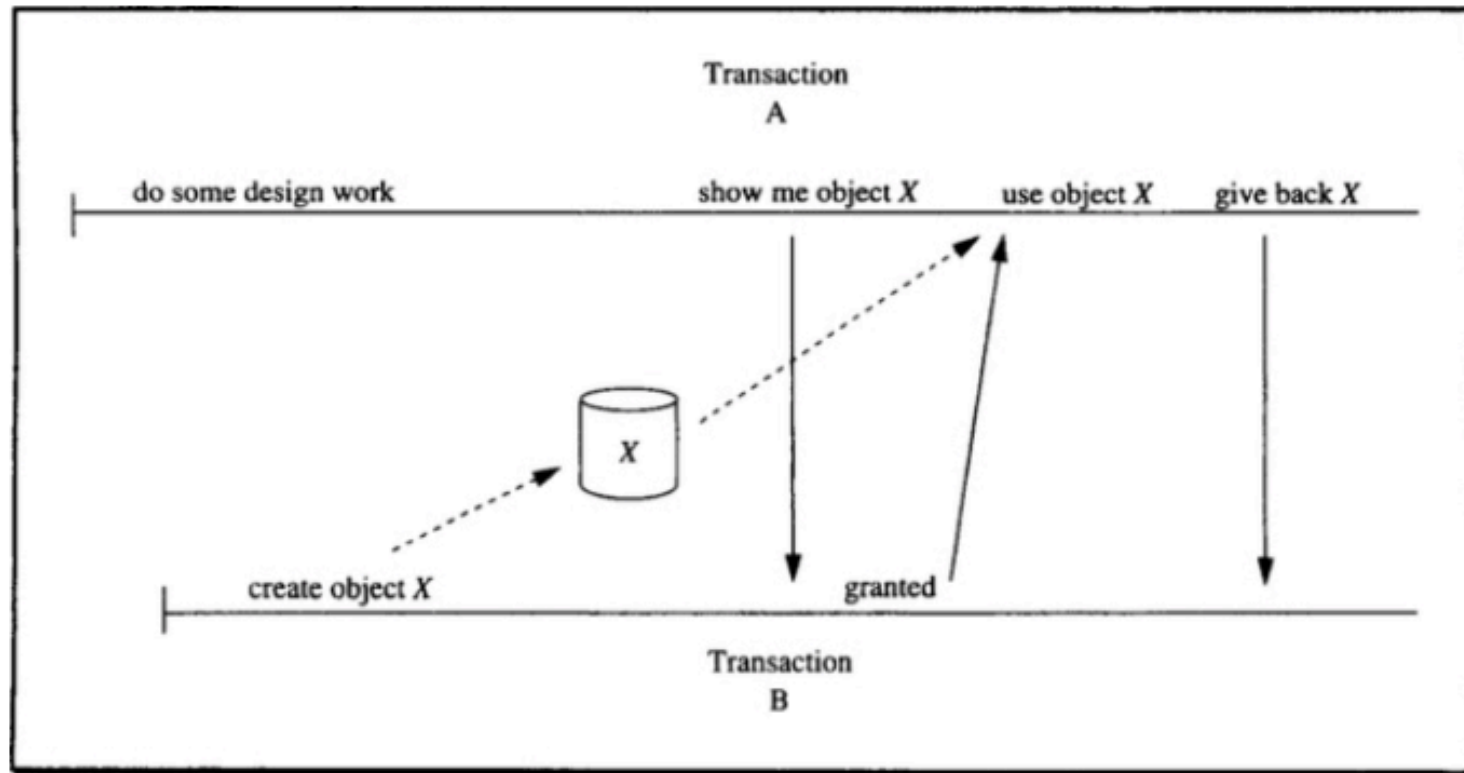


- Define chain of transactions as a unit of control
- Use compensation idea from Multi-level transaction

$S_1, S_2, \dots, S_{n-1}, S_n$

S_1, S_2, \dots, S_j (**abort**) , CS_{j-1}, CS_2, CS_1

Exotics



1. Prerelease upon request
2. Explicit return

Highly depend on Application Semantics